

## Importing libraries

In [0]:

```
# Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.py

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.layers import BatchNormalization

batch_size = 128
num_classes = 10
epochs = 15

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

Using TensorFlow backend.

```
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz (http://s3.amazonaws.com/img-datasets/mnist.npz)
11493376/11490434 [=====] - 2s 0us/step
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

In [0]:

```
# plotting train and test error

import matplotlib.pyplot as plt
%matplotlib inline

def plot_loss(training, epochs):
    fig, ax = plt.subplots(1,1)
    ax.set_xlabel('epoch')
    ax.set_ylabel('Categorical Crossentropy Loss')
    x_values = list(range(1,epochs+1))

    validation_loss = training.history['val_loss']
    train_loss = training.history['loss']

    ax.plot(x_values, validation_loss, 'b', label="Validation Loss")
    ax.plot(x_values, train_loss, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
```

**2 conv layers, kernels =(3x3), Max pooling =(2x2)**

In [0]:

```
%%time

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_5 (Conv2D)	(None, 26, 26, 32)	320
conv2d_6 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_3 (MaxPooling2	(None, 12, 12, 64)	0
dropout_5 (Dropout)	(None, 12, 12, 64)	0
flatten_3 (Flatten)	(None, 9216)	0
dense_5 (Dense)	(None, 128)	1179776
dropout_6 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 10)	1290
=====		
Total params: 1,199,882		
Trainable params: 1,199,882		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples


Epoch 1/15

60000/60000 [=====] - 147s 2ms/step - loss: 0.25

```
76 - accuracy: 0.9207 - val_loss: 0.0538 - val_accuracy: 0.9826
Epoch 2/15
60000/60000 [=====] - 145s 2ms/step - loss: 0.08
76 - accuracy: 0.9737 - val_loss: 0.0409 - val_accuracy: 0.9856
Epoch 3/15
60000/60000 [=====] - 144s 2ms/step - loss: 0.06
58 - accuracy: 0.9802 - val_loss: 0.0379 - val_accuracy: 0.9872
Epoch 4/15
60000/60000 [=====] - 144s 2ms/step - loss: 0.05
46 - accuracy: 0.9835 - val_loss: 0.0319 - val_accuracy: 0.9890
Epoch 5/15
60000/60000 [=====] - 145s 2ms/step - loss: 0.04
82 - accuracy: 0.9856 - val_loss: 0.0299 - val_accuracy: 0.9900
Epoch 6/15
60000/60000 [=====] - 147s 2ms/step - loss: 0.04
18 - accuracy: 0.9872 - val_loss: 0.0263 - val_accuracy: 0.9908
Epoch 7/15
60000/60000 [=====] - 147s 2ms/step - loss: 0.03
77 - accuracy: 0.9879 - val_loss: 0.0310 - val_accuracy: 0.9898
Epoch 8/15
60000/60000 [=====] - 148s 2ms/step - loss: 0.03
60 - accuracy: 0.9892 - val_loss: 0.0246 - val_accuracy: 0.9917
Epoch 9/15
60000/60000 [=====] - 148s 2ms/step - loss: 0.03
19 - accuracy: 0.9905 - val_loss: 0.0279 - val_accuracy: 0.9912
Epoch 10/15
60000/60000 [=====] - 148s 2ms/step - loss: 0.03
08 - accuracy: 0.9905 - val_loss: 0.0242 - val_accuracy: 0.9921
Epoch 11/15
60000/60000 [=====] - 151s 3ms/step - loss: 0.02
85 - accuracy: 0.9911 - val_loss: 0.0280 - val_accuracy: 0.9911
Epoch 12/15
60000/60000 [=====] - 149s 2ms/step - loss: 0.02
73 - accuracy: 0.9913 - val_loss: 0.0292 - val_accuracy: 0.9914
Epoch 13/15
60000/60000 [=====] - 145s 2ms/step - loss: 0.02
54 - accuracy: 0.9923 - val_loss: 0.0300 - val_accuracy: 0.9912
Epoch 14/15
60000/60000 [=====] - 145s 2ms/step - loss: 0.02
40 - accuracy: 0.9927 - val_loss: 0.0286 - val_accuracy: 0.9915
Epoch 15/15
60000/60000 [=====] - 150s 3ms/step - loss: 0.02
43 - accuracy: 0.9923 - val_loss: 0.0311 - val_accuracy: 0.9905
Test loss: 0.03108858480632407
Test accuracy: 0.9904999732971191
CPU times: user 1h 9min 25s, sys: 1min 15s, total: 1h 10min 40s
Wall time: 36min 50s
```

In [0]:

```
plot_loss(training, epochs)
```



from epoch 11 it starts overfitting. Best epoch is 10

**2 conv layers, kernels =(5x5) max pooling = (2X2)**

In [0]:

```

%%time
model = Sequential()
model.add(Conv2D(32, kernel_size=(5, 5),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_9 (Conv2D)	(None, 24, 24, 32)	832
conv2d_10 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_5 (MaxPooling2	(None, 10, 10, 64)	0
dropout_9 (Dropout)	(None, 10, 10, 64)	0
flatten_5 (Flatten)	(None, 6400)	0
dense_9 (Dense)	(None, 128)	819328
dropout_10 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 10)	1290
=====		
Total params: 872,714		
Trainable params: 872,714		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 232s 4ms/step - loss: 0.22

54 - accuracy: 0.9303 - val\_loss: 0.0528 - val\_accuracy: 0.9835

Epoch 2/15


60000/60000 [=====] - 230s 4ms/step - loss: 0.07

```
42 - accuracy: 0.9779 - val_loss: 0.0325 - val_accuracy: 0.9886
Epoch 3/15
60000/60000 [=====] - 229s 4ms/step - loss: 0.05
65 - accuracy: 0.9836 - val_loss: 0.0407 - val_accuracy: 0.9874
Epoch 4/15
60000/60000 [=====] - 228s 4ms/step - loss: 0.04
56 - accuracy: 0.9867 - val_loss: 0.0272 - val_accuracy: 0.9906
Epoch 5/15
60000/60000 [=====] - 229s 4ms/step - loss: 0.04
04 - accuracy: 0.9881 - val_loss: 0.0249 - val_accuracy: 0.9913
Epoch 6/15
60000/60000 [=====] - 229s 4ms/step - loss: 0.03
46 - accuracy: 0.9894 - val_loss: 0.0248 - val_accuracy: 0.9916
Epoch 7/15
60000/60000 [=====] - 234s 4ms/step - loss: 0.02
99 - accuracy: 0.9910 - val_loss: 0.0226 - val_accuracy: 0.9918
Epoch 8/15
60000/60000 [=====] - 230s 4ms/step - loss: 0.02
85 - accuracy: 0.9916 - val_loss: 0.0227 - val_accuracy: 0.9920
Epoch 9/15
60000/60000 [=====] - 230s 4ms/step - loss: 0.02
43 - accuracy: 0.9923 - val_loss: 0.0206 - val_accuracy: 0.9929
Epoch 10/15
60000/60000 [=====] - 236s 4ms/step - loss: 0.02
18 - accuracy: 0.9934 - val_loss: 0.0198 - val_accuracy: 0.9932
Epoch 11/15
60000/60000 [=====] - 230s 4ms/step - loss: 0.02
22 - accuracy: 0.9932 - val_loss: 0.0204 - val_accuracy: 0.9937
Epoch 12/15
60000/60000 [=====] - 235s 4ms/step - loss: 0.02
00 - accuracy: 0.9936 - val_loss: 0.0219 - val_accuracy: 0.9930
Epoch 13/15
60000/60000 [=====] - 237s 4ms/step - loss: 0.01
93 - accuracy: 0.9939 - val_loss: 0.0203 - val_accuracy: 0.9936
Epoch 14/15
60000/60000 [=====] - 237s 4ms/step - loss: 0.01
74 - accuracy: 0.9945 - val_loss: 0.0214 - val_accuracy: 0.9929
Epoch 15/15
60000/60000 [=====] - 236s 4ms/step - loss: 0.01
60 - accuracy: 0.9948 - val_loss: 0.0207 - val_accuracy: 0.9934
Test loss: 0.020748127230865065
Test accuracy: 0.993399977684021
CPU times: user 1h 50min, sys: 1min 20s, total: 1h 51min 21s
Wall time: 58min 12s
```



In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 11. It starts overfitting after 11 epochs. Here best accuracy is 99.37

**3 conv layers, kernels = (7x7),(5x5),(3x3), max pooling =(2x2)**



In [0]:

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(7, 7),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
=====		
conv2d_11 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_12 (Conv2D)	(None, 18, 18, 64)	51264
max_pooling2d_6 (MaxPooling2	(None, 9, 9, 64)	0
dropout_11 (Dropout)	(None, 9, 9, 64)	0
conv2d_13 (Conv2D)	(None, 7, 7, 128)	73856
max_pooling2d_7 (MaxPooling2	(None, 3, 3, 128)	0
dropout_12 (Dropout)	(None, 3, 3, 128)	0
flatten_6 (Flatten)	(None, 1152)	0
dense_11 (Dense)	(None, 512)	590336
dropout_13 (Dropout)	(None, 512)	0
dense_12 (Dense)	(None, 10)	5130
=====		

Total params: 722,186  
Trainable params: 722,186  
Non-trainable params: 0

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.2804 - accuracy: 0.9097 - val\_loss: 0.0493 - val\_accuracy: 0.9831

Epoch 2/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.0775 - accuracy: 0.9762 - val\_loss: 0.0334 - val\_accuracy: 0.9890

Epoch 3/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.0542 - accuracy: 0.9834 - val\_loss: 0.0306 - val\_accuracy: 0.9898

Epoch 4/15

60000/60000 [=====] - 226s 4ms/step - loss: 0.0456 - accuracy: 0.9862 - val\_loss: 0.0238 - val\_accuracy: 0.9922

Epoch 5/15

60000/60000 [=====] - 227s 4ms/step - loss: 0.0384 - accuracy: 0.9882 - val\_loss: 0.0206 - val\_accuracy: 0.9932

Epoch 6/15

60000/60000 [=====] - 230s 4ms/step - loss: 0.0331 - accuracy: 0.9899 - val\_loss: 0.0193 - val\_accuracy: 0.9941

Epoch 7/15

60000/60000 [=====] - 230s 4ms/step - loss: 0.0308 - accuracy: 0.9904 - val\_loss: 0.0190 - val\_accuracy: 0.9939

Epoch 8/15

60000/60000 [=====] - 230s 4ms/step - loss: 0.0281 - accuracy: 0.9912 - val\_loss: 0.0234 - val\_accuracy: 0.9922

Epoch 9/15

60000/60000 [=====] - 237s 4ms/step - loss: 0.0253 - accuracy: 0.9922 - val\_loss: 0.0195 - val\_accuracy: 0.9941

Epoch 10/15

60000/60000 [=====] - 230s 4ms/step - loss: 0.0229 - accuracy: 0.9928 - val\_loss: 0.0170 - val\_accuracy: 0.9939

Epoch 11/15

60000/60000 [=====] - 227s 4ms/step - loss: 0.0224 - accuracy: 0.9926 - val\_loss: 0.0149 - val\_accuracy: 0.9953

Epoch 12/15

60000/60000 [=====] - 234s 4ms/step - loss: 0.0192 - accuracy: 0.9938 - val\_loss: 0.0168 - val\_accuracy: 0.9950

Epoch 13/15

60000/60000 [=====] - 229s 4ms/step - loss: 0.0185 - accuracy: 0.9941 - val\_loss: 0.0150 - val\_accuracy: 0.9955

Epoch 14/15

60000/60000 [=====] - 230s 4ms/step - loss: 0.0173 - accuracy: 0.9948 - val\_loss: 0.0187 - val\_accuracy: 0.9941

Epoch 15/15

60000/60000 [=====] - 233s 4ms/step - loss: 0.0169 - accuracy: 0.9947 - val\_loss: 0.0188 - val\_accuracy: 0.9951


Test loss: 0.01883504462489873

Test accuracy: 0.9951000213623047



In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 13. Its overfitting from epoch 14 onwards. Best accuracy is 99.55%

**2 conv layers, kernels =(3x3),(3x3),Max pooling = (2x2), padding = same**

In [0]:

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu', padding='same',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
conv2d_14 (Conv2D)	(None, 28, 28, 32)	320
conv2d_15 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_14 (Dropout)	(None, 14, 14, 64)	0
flatten_7 (Flatten)	(None, 12544)	0
dense_13 (Dense)	(None, 128)	1605760
dropout_15 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
Total params: 1,625,866		
Trainable params: 1,625,866		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 193s 3ms/step - loss: 0.2618  
 - accuracy: 0.9196 - val\_loss: 0.0563 - val\_accuracy: 0.9819

Epoch 2/15


60000/60000 [=====] - 190s 3ms/step - loss: 0.0897  
 - accuracy: 0.9735 - val\_loss: 0.0394 - val\_accuracy: 0.9866

Epoch 3/15

```
60000/60000 [=====] - 190s 3ms/step - loss: 0.0656
- accuracy: 0.9803 - val_loss: 0.0325 - val_accuracy: 0.9885
Epoch 4/15
60000/60000 [=====] - 190s 3ms/step - loss: 0.0554
- accuracy: 0.9833 - val_loss: 0.0310 - val_accuracy: 0.9894
Epoch 5/15
60000/60000 [=====] - 190s 3ms/step - loss: 0.0471
- accuracy: 0.9860 - val_loss: 0.0295 - val_accuracy: 0.9900
Epoch 6/15
60000/60000 [=====] - 194s 3ms/step - loss: 0.0414
- accuracy: 0.9875 - val_loss: 0.0346 - val_accuracy: 0.9887
Epoch 7/15
60000/60000 [=====] - 194s 3ms/step - loss: 0.0382
- accuracy: 0.9885 - val_loss: 0.0272 - val_accuracy: 0.9905
Epoch 8/15
60000/60000 [=====] - 197s 3ms/step - loss: 0.0358
- accuracy: 0.9890 - val_loss: 0.0281 - val_accuracy: 0.9910
Epoch 9/15
60000/60000 [=====] - 199s 3ms/step - loss: 0.0338
- accuracy: 0.9897 - val_loss: 0.0286 - val_accuracy: 0.9912
Epoch 10/15
60000/60000 [=====] - 203s 3ms/step - loss: 0.0318
- accuracy: 0.9903 - val_loss: 0.0279 - val_accuracy: 0.9917
Epoch 11/15
60000/60000 [=====] - 210s 3ms/step - loss: 0.0295
- accuracy: 0.9912 - val_loss: 0.0298 - val_accuracy: 0.9908
Epoch 12/15
60000/60000 [=====] - 206s 3ms/step - loss: 0.0283
- accuracy: 0.9917 - val_loss: 0.0272 - val_accuracy: 0.9921
Epoch 13/15
60000/60000 [=====] - 197s 3ms/step - loss: 0.0260
- accuracy: 0.9918 - val_loss: 0.0303 - val_accuracy: 0.9904
Epoch 14/15
60000/60000 [=====] - 200s 3ms/step - loss: 0.0254
- accuracy: 0.9919 - val_loss: 0.0319 - val_accuracy: 0.9920
Epoch 15/15
60000/60000 [=====] - 196s 3ms/step - loss: 0.0256
- accuracy: 0.9923 - val_loss: 0.0287 - val_accuracy: 0.9917
Test loss: 0.028727352098264054
Test accuracy: 0.9916999936103821
```

In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 10. Its overfitting from epoch 11. Best accuracy is 99.17%

**3 conv layers, kernels = (7x7),(5x5),(3x3), padding =same**

In [0]:

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(7, 7),
                 activation='relu', padding='same',
                 input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3,3), activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.3))
model.add(Dense(128, activation='relu'))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
=====		
conv2d_16 (Conv2D)	(None, 28, 28, 32)	1600
conv2d_17 (Conv2D)	(None, 28, 28, 64)	51264
max_pooling2d_9 (MaxPooling2	(None, 14, 14, 64)	0
dropout_16 (Dropout)	(None, 14, 14, 64)	0
conv2d_18 (Conv2D)	(None, 14, 14, 128)	73856
max_pooling2d_10 (MaxPooling	(None, 7, 7, 128)	0
dropout_17 (Dropout)	(None, 7, 7, 128)	0
dense_15 (Dense)	(None, 7, 7, 128)	16512
flatten_8 (Flatten)	(None, 6272)	0
dense_16 (Dense)	(None, 128)	802944
dropout_18 (Dropout)	(None, 128)	0

dense\_17 (Dense) (None, 10) 1290

=====

Total params: 947,466

Trainable params: 947,466

Non-trainable params: 0

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 568s 9ms/step - loss: 0.28

78 - accuracy: 0.9090 - val\_loss: 0.0494 - val\_accuracy: 0.9842

Epoch 2/15

60000/60000 [=====] - 575s 10ms/step - loss: 0.0

841 - accuracy: 0.9751 - val\_loss: 0.0300 - val\_accuracy: 0.9901

Epoch 3/15

60000/60000 [=====] - 573s 10ms/step - loss: 0.0

623 - accuracy: 0.9818 - val\_loss: 0.0283 - val\_accuracy: 0.9905

Epoch 4/15

60000/60000 [=====] - 563s 9ms/step - loss: 0.05

16 - accuracy: 0.9850 - val\_loss: 0.0251 - val\_accuracy: 0.9915

Epoch 5/15

60000/60000 [=====] - 561s 9ms/step - loss: 0.04

36 - accuracy: 0.9873 - val\_loss: 0.0270 - val\_accuracy: 0.9907

Epoch 6/15

60000/60000 [=====] - 567s 9ms/step - loss: 0.03

80 - accuracy: 0.9891 - val\_loss: 0.0184 - val\_accuracy: 0.9934

Epoch 7/15

60000/60000 [=====] - 567s 9ms/step - loss: 0.03

50 - accuracy: 0.9897 - val\_loss: 0.0172 - val\_accuracy: 0.9943

Epoch 8/15

60000/60000 [=====] - 571s 10ms/step - loss: 0.0

311 - accuracy: 0.9904 - val\_loss: 0.0180 - val\_accuracy: 0.9940

Epoch 9/15

60000/60000 [=====] - 582s 10ms/step - loss: 0.0

293 - accuracy: 0.9913 - val\_loss: 0.0187 - val\_accuracy: 0.9936

Epoch 10/15

60000/60000 [=====] - 635s 11ms/step - loss: 0.0

257 - accuracy: 0.9924 - val\_loss: 0.0197 - val\_accuracy: 0.9940

Epoch 11/15

60000/60000 [=====] - 596s 10ms/step - loss: 0.0

244 - accuracy: 0.9926 - val\_loss: 0.0156 - val\_accuracy: 0.9948

Epoch 12/15

60000/60000 [=====] - 569s 9ms/step - loss: 0.02

38 - accuracy: 0.9930 - val\_loss: 0.0167 - val\_accuracy: 0.9940

Epoch 13/15

60000/60000 [=====] - 573s 10ms/step - loss: 0.0

217 - accuracy: 0.9937 - val\_loss: 0.0180 - val\_accuracy: 0.9942

Epoch 14/15

60000/60000 [=====] - 577s 10ms/step - loss: 0.0

211 - accuracy: 0.9938 - val\_loss: 0.0156 - val\_accuracy: 0.9946

Epoch 15/15

60000/60000 [=====] - 591s 10ms/step - loss: 0.0

188 - accuracy: 0.9946 - val\_loss: 0.0199 - val\_accuracy: 0.9931

Test loss: 0.01988242387709852


Test accuracy: 0.9930999875068665





In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 7. Best accuracy 99.43%

**3 conv layers, kernels=(5x5), padding = same, strides= 1**

In [0]:

```

%%time
model = Sequential()
model.add(Conv2D(32, kernel_size=(5, 5),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5),strides=1, activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(80, (5,5), activation= 'relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding = 'same',strides=1))
model.add(Dropout(0.75))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
=====		
conv2d_19 (Conv2D)	(None, 24, 24, 32)	832
conv2d_20 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_11 (MaxPooling)	(None, 10, 10, 64)	0
dropout_19 (Dropout)	(None, 10, 10, 64)	0
conv2d_21 (Conv2D)	(None, 6, 6, 80)	128080
max_pooling2d_12 (MaxPooling)	(None, 6, 6, 80)	0
dropout_20 (Dropout)	(None, 6, 6, 80)	0
flatten_9 (Flatten)	(None, 2880)	0
dense_18 (Dense)	(None, 128)	368768
dropout_21 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
=====		

Total params: 550,234  
Trainable params: 550,234  
Non-trainable params: 0

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 286s 5ms/step - loss: 0.3174  
- accuracy: 0.8973 - val\_loss: 0.0386 - val\_accuracy: 0.9877

Epoch 2/15

60000/60000 [=====] - 281s 5ms/step - loss: 0.0798  
- accuracy: 0.9762 - val\_loss: 0.0317 - val\_accuracy: 0.9899

Epoch 3/15

60000/60000 [=====] - 275s 5ms/step - loss: 0.0622  
- accuracy: 0.9821 - val\_loss: 0.0224 - val\_accuracy: 0.9926

Epoch 4/15

60000/60000 [=====] - 275s 5ms/step - loss: 0.0500  
- accuracy: 0.9855 - val\_loss: 0.0240 - val\_accuracy: 0.9919

Epoch 5/15

60000/60000 [=====] - 276s 5ms/step - loss: 0.0467  
- accuracy: 0.9860 - val\_loss: 0.0204 - val\_accuracy: 0.9935

Epoch 6/15

60000/60000 [=====] - 276s 5ms/step - loss: 0.0399  
- accuracy: 0.9884 - val\_loss: 0.0186 - val\_accuracy: 0.9941

Epoch 7/15

60000/60000 [=====] - 277s 5ms/step - loss: 0.0383  
- accuracy: 0.9891 - val\_loss: 0.0189 - val\_accuracy: 0.9943

Epoch 8/15

60000/60000 [=====] - 276s 5ms/step - loss: 0.0357  
- accuracy: 0.9895 - val\_loss: 0.0178 - val\_accuracy: 0.9939

Epoch 9/15

60000/60000 [=====] - 275s 5ms/step - loss: 0.0335  
- accuracy: 0.9904 - val\_loss: 0.0183 - val\_accuracy: 0.9942

Epoch 10/15

60000/60000 [=====] - 281s 5ms/step - loss: 0.0314  
- accuracy: 0.9908 - val\_loss: 0.0202 - val\_accuracy: 0.9940

Epoch 11/15

60000/60000 [=====] - 276s 5ms/step - loss: 0.0293  
- accuracy: 0.9911 - val\_loss: 0.0255 - val\_accuracy: 0.9930

Epoch 12/15

60000/60000 [=====] - 283s 5ms/step - loss: 0.0309  
- accuracy: 0.9909 - val\_loss: 0.0212 - val\_accuracy: 0.9929

Epoch 13/15

60000/60000 [=====] - 276s 5ms/step - loss: 0.0276  
- accuracy: 0.9919 - val\_loss: 0.0182 - val\_accuracy: 0.9944

Epoch 14/15

60000/60000 [=====] - 276s 5ms/step - loss: 0.0260  
- accuracy: 0.9925 - val\_loss: 0.0213 - val\_accuracy: 0.9941

Epoch 15/15

60000/60000 [=====] - 279s 5ms/step - loss: 0.0257  
- accuracy: 0.9931 - val\_loss: 0.0183 - val\_accuracy: 0.9944

Test loss: 0.018290850276246533


Test accuracy: 0.9944000244140625

CPU times: user 2h 11min 45s, sys: 1min 44s, total: 2h 13min 29s

Wall time: 1h 9min 40s

In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 9. Best accuracy 99.42%

**3 conv layers, kernels =(7x7), padding = same, strides =1,2**

In [0]:

```

%%time
model = Sequential()
model.add(Conv2D(32, kernel_size=(7, 7),
                activation='relu', padding='same', strides=1,
                input_shape=input_shape))
model.add(Conv2D(64, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(80, (7,7), activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same',strides=2))
model.add(Dropout(0.75))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 28, 28, 32)	1600
conv2d_2 (Conv2D)	(None, 22, 22, 64)	100416
max_pooling2d_1 (MaxPooling2D)	(None, 11, 11, 64)	0
dropout_1 (Dropout)	(None, 11, 11, 64)	0
conv2d_3 (Conv2D)	(None, 5, 5, 80)	250960
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 80)	0
dropout_2 (Dropout)	(None, 3, 3, 80)	0
flatten_1 (Flatten)	(None, 720)	0
dense_1 (Dense)	(None, 128)	92288
dropout_3 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290
=====		

Total params: 446,554  
Trainable params: 446,554  
Non-trainable params: 0

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 568s 9ms/step - loss: 0.4405 - accuracy: 0.8557 - val\_loss: 0.0462 - val\_accuracy: 0.9856

Epoch 2/15

60000/60000 [=====] - 566s 9ms/step - loss: 0.1057 - accuracy: 0.9701 - val\_loss: 0.0371 - val\_accuracy: 0.9880

Epoch 3/15

60000/60000 [=====] - 566s 9ms/step - loss: 0.0766 - accuracy: 0.9791 - val\_loss: 0.0285 - val\_accuracy: 0.9914

Epoch 4/15

60000/60000 [=====] - 572s 10ms/step - loss: 0.0654 - accuracy: 0.9825 - val\_loss: 0.0217 - val\_accuracy: 0.9932

Epoch 5/15

60000/60000 [=====] - 573s 10ms/step - loss: 0.0564 - accuracy: 0.9847 - val\_loss: 0.0222 - val\_accuracy: 0.9931

Epoch 6/15

60000/60000 [=====] - 568s 9ms/step - loss: 0.0516 - accuracy: 0.9860 - val\_loss: 0.0228 - val\_accuracy: 0.9927

Epoch 7/15

60000/60000 [=====] - 566s 9ms/step - loss: 0.0492 - accuracy: 0.9868 - val\_loss: 0.0215 - val\_accuracy: 0.9933

Epoch 8/15

60000/60000 [=====] - 566s 9ms/step - loss: 0.0441 - accuracy: 0.9880 - val\_loss: 0.0206 - val\_accuracy: 0.9932

Epoch 9/15

60000/60000 [=====] - 569s 9ms/step - loss: 0.0422 - accuracy: 0.9885 - val\_loss: 0.0195 - val\_accuracy: 0.9942

Epoch 10/15

60000/60000 [=====] - 574s 10ms/step - loss: 0.0391 - accuracy: 0.9886 - val\_loss: 0.0183 - val\_accuracy: 0.9942

Epoch 11/15

60000/60000 [=====] - 574s 10ms/step - loss: 0.0356 - accuracy: 0.9902 - val\_loss: 0.0202 - val\_accuracy: 0.9937

Epoch 12/15

60000/60000 [=====] - 566s 9ms/step - loss: 0.0330 - accuracy: 0.9909 - val\_loss: 0.0222 - val\_accuracy: 0.9938

Epoch 13/15

60000/60000 [=====] - 564s 9ms/step - loss: 0.0357 - accuracy: 0.9895 - val\_loss: 0.0195 - val\_accuracy: 0.9940

Epoch 14/15

60000/60000 [=====] - 567s 9ms/step - loss: 0.0326 - accuracy: 0.9905 - val\_loss: 0.0228 - val\_accuracy: 0.9939

Epoch 15/15

60000/60000 [=====] - 566s 9ms/step - loss: 0.0326 - accuracy: 0.9908 - val\_loss: 0.0220 - val\_accuracy: 0.9930

Test loss: 0.02198781901857328

Test accuracy: 0.9929999709129333


CPU times: user 4h 33min 20s, sys: 3min 5s, total: 4h 36min 25s

Wall time: 2h 22min 26s



In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 5. best accuracy 99.31%

**3 conv layers, kernels = (7x7), max pooling =(2x2),  
strides=2,1**

In [0]:

```

%%time
model = Sequential()
model.add(Conv2D(32, kernel_size=(7, 7),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),strides=2))
model.add(BatchNormalization())
model.add(Dropout(0.25))

model.add(Conv2D(80, (7,7), activation= 'relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding = 'same',strides=1))
model.add(BatchNormalization())
model.add(Dropout(0.75))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_5 (Conv2D)	(None, 16, 16, 64)	100416
max_pooling2d_3 (MaxPooling2	(None, 8, 8, 64)	0
batch_normalization_1 (Batch	(None, 8, 8, 64)	256
dropout_4 (Dropout)	(None, 8, 8, 64)	0
conv2d_6 (Conv2D)	(None, 2, 2, 80)	250960
max_pooling2d_4 (MaxPooling2	(None, 2, 2, 80)	0
batch_normalization_2 (Batch	(None, 2, 2, 80)	320
dropout_5 (Dropout)	(None, 2, 2, 80)	0
flatten_2 (Flatten)	(None, 320)	0



dense_3 (Dense)	(None, 128)	41088
-----------------	-------------	-------

dropout_6 (Dropout)	(None, 128)	0
---------------------	-------------	---

dense_4 (Dense)	(None, 10)	1290
-----------------	------------	------

=====

Total params: 395,930

Trainable params: 395,642

Non-trainable params: 288

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 323s 5ms/step - loss: 0.4930

- accuracy: 0.8647 - val\_loss: 0.0750 - val\_accuracy: 0.9799

Epoch 2/15

60000/60000 [=====] - 318s 5ms/step - loss: 0.1116

- accuracy: 0.9691 - val\_loss: 0.0356 - val\_accuracy: 0.9884

Epoch 3/15

60000/60000 [=====] - 320s 5ms/step - loss: 0.0749

- accuracy: 0.9794 - val\_loss: 0.0318 - val\_accuracy: 0.9908

Epoch 4/15

60000/60000 [=====] - 316s 5ms/step - loss: 0.0630

- accuracy: 0.9823 - val\_loss: 0.0301 - val\_accuracy: 0.9914

Epoch 5/15

60000/60000 [=====] - 315s 5ms/step - loss: 0.0585

- accuracy: 0.9840 - val\_loss: 0.0242 - val\_accuracy: 0.9926

Epoch 6/15

60000/60000 [=====] - 316s 5ms/step - loss: 0.0482

- accuracy: 0.9866 - val\_loss: 0.0243 - val\_accuracy: 0.9931

Epoch 7/15

60000/60000 [=====] - 317s 5ms/step - loss: 0.0407

- accuracy: 0.9885 - val\_loss: 0.0225 - val\_accuracy: 0.9934

Epoch 8/15

60000/60000 [=====] - 317s 5ms/step - loss: 0.0387

- accuracy: 0.9891 - val\_loss: 0.0298 - val\_accuracy: 0.9920

Epoch 9/15

60000/60000 [=====] - 318s 5ms/step - loss: 0.0364

- accuracy: 0.9897 - val\_loss: 0.0191 - val\_accuracy: 0.9939

Epoch 10/15

60000/60000 [=====] - 318s 5ms/step - loss: 0.0340

- accuracy: 0.9905 - val\_loss: 0.0239 - val\_accuracy: 0.9931

Epoch 11/15

60000/60000 [=====] - 318s 5ms/step - loss: 0.0300

- accuracy: 0.9914 - val\_loss: 0.0390 - val\_accuracy: 0.9888

Epoch 12/15

60000/60000 [=====] - 323s 5ms/step - loss: 0.0297

- accuracy: 0.9918 - val\_loss: 0.0316 - val\_accuracy: 0.9919

Epoch 13/15

60000/60000 [=====] - 317s 5ms/step - loss: 0.0295

- accuracy: 0.9919 - val\_loss: 0.0229 - val\_accuracy: 0.9920

Epoch 14/15

60000/60000 [=====] - 320s 5ms/step - loss: 0.0274

- accuracy: 0.9924 - val\_loss: 0.0245 - val\_accuracy: 0.9915

Epoch 15/15

60000/60000 [=====] - 312s 5ms/step - loss: 0.0246

- accuracy: 0.9930 - val\_loss: 0.0230 - val\_accuracy: 0.9934

Test loss: 0.023033796002792678


Test accuracy: 0.993399977684021

CPU times: user 2h 26min 53s, sys: 2min 41s, total: 2h 29min 35s

Wall time: 1h 19min 39s

In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 7. Best accuracy 99.34%

**3 conv layers, kernels = (5x5), strides =1, max pooling = (2x2), padding = same**

In [0]:

```
%%time

model = Sequential()
model.add(Conv2D(32, kernel_size=(5, 5),
                activation='relu', strides =1,
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5),strides=1, activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.25))

model.add(Conv2D(80, (5,5), activation= 'relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding = 'same',strides=1))
model.add(BatchNormalization())
model.add(Dropout(0.75))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 24, 24, 32)	832
conv2d_8 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_5 (MaxPooling2	(None, 10, 10, 64)	0
batch_normalization_3 (Batch	(None, 10, 10, 64)	256
dropout_7 (Dropout)	(None, 10, 10, 64)	0
conv2d_9 (Conv2D)	(None, 6, 6, 80)	128080
max_pooling2d_6 (MaxPooling2	(None, 6, 6, 80)	0
batch_normalization_4 (Batch	(None, 6, 6, 80)	320
dropout_8 (Dropout)	(None, 6, 6, 80)	0
flatten_3 (Flatten)	(None, 2880)	0

dense_5 (Dense)	(None, 128)	368768
dropout_9 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 10)	1290
=====		
Total params: 550,810		
Trainable params: 550,522		
Non-trainable params: 288		

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 292s 5ms/step - loss: 0.3540  
- accuracy: 0.9021 - val\_loss: 0.2448 - val\_accuracy: 0.9415

Epoch 2/15

60000/60000 [=====] - 291s 5ms/step - loss: 0.1071  
- accuracy: 0.9697 - val\_loss: 0.0404 - val\_accuracy: 0.9883

Epoch 3/15

60000/60000 [=====] - 290s 5ms/step - loss: 0.0823  
- accuracy: 0.9772 - val\_loss: 0.0255 - val\_accuracy: 0.9911

Epoch 4/15

60000/60000 [=====] - 288s 5ms/step - loss: 0.0699  
- accuracy: 0.9796 - val\_loss: 0.0366 - val\_accuracy: 0.9885

Epoch 5/15

60000/60000 [=====] - 288s 5ms/step - loss: 0.0613  
- accuracy: 0.9823 - val\_loss: 0.0364 - val\_accuracy: 0.9886

Epoch 6/15

60000/60000 [=====] - 288s 5ms/step - loss: 0.0561  
- accuracy: 0.9842 - val\_loss: 0.0276 - val\_accuracy: 0.9921

Epoch 7/15

60000/60000 [=====] - 294s 5ms/step - loss: 0.0512  
- accuracy: 0.9853 - val\_loss: 0.0356 - val\_accuracy: 0.9900

Epoch 8/15

60000/60000 [=====] - 289s 5ms/step - loss: 0.0472  
- accuracy: 0.9869 - val\_loss: 0.0209 - val\_accuracy: 0.9931

Epoch 9/15

60000/60000 [=====] - 293s 5ms/step - loss: 0.0467  
- accuracy: 0.9873 - val\_loss: 0.0194 - val\_accuracy: 0.9940

Epoch 10/15

60000/60000 [=====] - 287s 5ms/step - loss: 0.0436  
- accuracy: 0.9877 - val\_loss: 0.0295 - val\_accuracy: 0.9917

Epoch 11/15

60000/60000 [=====] - 291s 5ms/step - loss: 0.0429  
- accuracy: 0.9876 - val\_loss: 0.0171 - val\_accuracy: 0.9948

Epoch 12/15

60000/60000 [=====] - 287s 5ms/step - loss: 0.0398  
- accuracy: 0.9884 - val\_loss: 0.0234 - val\_accuracy: 0.9942

Epoch 13/15

60000/60000 [=====] - 294s 5ms/step - loss: 0.0358  
- accuracy: 0.9901 - val\_loss: 0.0196 - val\_accuracy: 0.9943

Epoch 14/15

60000/60000 [=====] - 289s 5ms/step - loss: 0.0355  
- accuracy: 0.9900 - val\_loss: 0.0247 - val\_accuracy: 0.9928

Epoch 15/15

60000/60000 [=====] - 293s 5ms/step - loss: 0.0344  
- accuracy: 0.9903 - val\_loss: 0.0252 - val\_accuracy: 0.9930

Test loss: 0.02517386771799874


Test accuracy: 0.9929999709129333

CPU times: user 2h 16min 3s, sys: 2min 40s, total: 2h 18min 43s

Wall time: 1h 12min 44s

In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 9. best accuracy 99.40%

**3 conv layers, kernels=(3x3), max pooling =(2x2), padding =same, strides = 1**

In [0]:

```
%%time

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3),strides=1, activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.25))

model.add(Conv2D(80, (3,3), activation= 'relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding = 'same',strides=1))
model.add(BatchNormalization())
model.add(Dropout(0.75))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adam(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_10 (Conv2D)	(None, 26, 26, 32)	320
conv2d_11 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_7 (MaxPooling2	(None, 12, 12, 64)	0
batch_normalization_5 (Batch	(None, 12, 12, 64)	256
dropout_10 (Dropout)	(None, 12, 12, 64)	0
conv2d_12 (Conv2D)	(None, 10, 10, 80)	46160
max_pooling2d_8 (MaxPooling2	(None, 10, 10, 80)	0
batch_normalization_6 (Batch	(None, 10, 10, 80)	320
dropout_11 (Dropout)	(None, 10, 10, 80)	0
flatten_4 (Flatten)	(None, 8000)	0

dense_7 (Dense)	(None, 128)	1024128
dropout_12 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
=====		
Total params: 1,090,970		
Trainable params: 1,090,682		
Non-trainable params: 288		

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 209s 3ms/step - loss: 0.3466  
- accuracy: 0.9030 - val\_loss: 0.1810 - val\_accuracy: 0.9499

Epoch 2/15

60000/60000 [=====] - 209s 3ms/step - loss: 0.1333  
- accuracy: 0.9634 - val\_loss: 0.0575 - val\_accuracy: 0.9825

Epoch 3/15

60000/60000 [=====] - 210s 3ms/step - loss: 0.1070  
- accuracy: 0.9704 - val\_loss: 0.0333 - val\_accuracy: 0.9899

Epoch 4/15

60000/60000 [=====] - 211s 4ms/step - loss: 0.0898  
- accuracy: 0.9760 - val\_loss: 0.0327 - val\_accuracy: 0.9907

Epoch 5/15

60000/60000 [=====] - 211s 4ms/step - loss: 0.0761  
- accuracy: 0.9781 - val\_loss: 0.0310 - val\_accuracy: 0.9913

Epoch 6/15

60000/60000 [=====] - 215s 4ms/step - loss: 0.0718  
- accuracy: 0.9803 - val\_loss: 0.0253 - val\_accuracy: 0.9932

Epoch 7/15

60000/60000 [=====] - 209s 3ms/step - loss: 0.0654  
- accuracy: 0.9818 - val\_loss: 0.0380 - val\_accuracy: 0.9897

Epoch 8/15

60000/60000 [=====] - 209s 3ms/step - loss: 0.0631  
- accuracy: 0.9822 - val\_loss: 0.0234 - val\_accuracy: 0.9937

Epoch 9/15

60000/60000 [=====] - 214s 4ms/step - loss: 0.0556  
- accuracy: 0.9833 - val\_loss: 0.0250 - val\_accuracy: 0.9928

Epoch 10/15

60000/60000 [=====] - 208s 3ms/step - loss: 0.0547  
- accuracy: 0.9844 - val\_loss: 0.0300 - val\_accuracy: 0.9927

Epoch 11/15

60000/60000 [=====] - 209s 3ms/step - loss: 0.0521  
- accuracy: 0.9851 - val\_loss: 0.0283 - val\_accuracy: 0.9928

Epoch 12/15

60000/60000 [=====] - 210s 3ms/step - loss: 0.0518  
- accuracy: 0.9856 - val\_loss: 0.0289 - val\_accuracy: 0.9920

Epoch 13/15

60000/60000 [=====] - 209s 3ms/step - loss: 0.0487  
- accuracy: 0.9862 - val\_loss: 0.0211 - val\_accuracy: 0.9932

Epoch 14/15

60000/60000 [=====] - 210s 3ms/step - loss: 0.0496  
- accuracy: 0.9866 - val\_loss: 0.0218 - val\_accuracy: 0.9930

Epoch 15/15

60000/60000 [=====] - 212s 4ms/step - loss: 0.0443  
- accuracy: 0.9871 - val\_loss: 0.0225 - val\_accuracy: 0.9939

Test loss: 0.022474063225078576


Test accuracy: 0.9939000010490417

CPU times: user 1h 39min 13s, sys: 2min 8s, total: 1h 41min 21s

Wall time: 52min 45s

In [0]:

```
# plotting train and test error
plot_loss(training, epochs)
```



best epoch is 8. best accuracy 99.37%

## Summary

In [1]:

```
from prettytable import PrettyTable
pt = PrettyTable()
pt.field_names = ["Layers", "Kernels", "Test Accuracy", "epoch"]
pt.add_row(["2", "(3x3), (3x3)", "99.21", "10"])
pt.add_row(["2", "(5x5), (5x5)", "99.37", "11"])
pt.add_row(["3", "(7x7), (5x5), (3x3)", "99.55", "13"])
print(pt)
```

Layers	Kernels	Test Accuracy	epoch
2	(3x3), (3x3)	99.21	10
2	(5x5), (5x5)	99.37	11
3	(7x7), (5x5), (3x3)	99.55	13



In [2]:

```
pt = PrettyTable()
pt.field_names = ["Layers", "Kernels", "Test Accuracy", "epoch", "padding"]
pt.add_row(["2", "(3x3),(3x3)", "99.17", "10", "same"])
pt.add_row(["3", "(7x7),(5x5),(3x3)", "99.43", "7", "same"])
print(pt)
```

Layers	Kernels	Test Accuracy	epoch	padding
2	(3x3),(3x3)	99.17	10	same
3	(7x7),(5x5),(3x3)	99.43	7	same

In [3]:

```
pt = PrettyTable()
pt.field_names = ["Layers", "Kernels", "Test Accuracy", "epoch", "padding", "strides"]
pt.add_row(["3", "(3x3),(3x3),(3x3)", "99.42", "9", "same", "1"])
pt.add_row(["3", "(5x5),(5x5),(5x5)", "99.34", "5", "same", "1"])
pt.add_row(["3", "(7x7),(7x7),(7x7)", "99.40", "9", "same", "2"])
pt.add_row(["3", "(7x7),(7x7),(7x7)", "99.37", "8", "same", "1,2"])
print(pt)
```

Layers	Kernels	Test Accuracy	epoch	padding	strides
3	(3x3),(3x3),(3x3)	99.42	9	same	1
3	(5x5),(5x5),(5x5)	99.34	5	same	1
3	(7x7),(7x7),(7x7)	99.40	9	same	2
3	(7x7),(7x7),(7x7)	99.37	8	same	1,2

## 3 conv layers, kernels = (7x7),(5x5),(3x3), Using sigmoid

In [0]:

```
%%time

model = Sequential()
model.add(Conv2D(32, kernel_size=(7, 7),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='sigmoid'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3,3), activation='sigmoid'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(512, activation='sigmoid'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_8 (Conv2D)	(None, 18, 18, 64)	51264
max_pooling2d_5 (MaxPooling2	(None, 9, 9, 64)	0
dropout_7 (Dropout)	(None, 9, 9, 64)	0
conv2d_9 (Conv2D)	(None, 7, 7, 128)	73856
max_pooling2d_6 (MaxPooling2	(None, 3, 3, 128)	0
dropout_8 (Dropout)	(None, 3, 3, 128)	0
flatten_3 (Flatten)	(None, 1152)	0
dense_5 (Dense)	(None, 512)	590336
dropout_9 (Dropout)	(None, 512)	0

dense\_6 (Dense) (None, 10) 5130

=====

Total params: 722,186

Trainable params: 722,186

Non-trainable params: 0

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 226s 4ms/step - loss: 2.02

95 - accuracy: 0.2398 - val\_loss: 0.4800 - val\_accuracy: 0.8586

Epoch 2/15

60000/60000 [=====] - 226s 4ms/step - loss: 0.30

56 - accuracy: 0.9069 - val\_loss: 0.1415 - val\_accuracy: 0.9534

Epoch 3/15

60000/60000 [=====] - 227s 4ms/step - loss: 0.16

39 - accuracy: 0.9489 - val\_loss: 0.0888 - val\_accuracy: 0.9737

Epoch 4/15

60000/60000 [=====] - 227s 4ms/step - loss: 0.12

28 - accuracy: 0.9620 - val\_loss: 0.0801 - val\_accuracy: 0.9747

Epoch 5/15

60000/60000 [=====] - 226s 4ms/step - loss: 0.10

18 - accuracy: 0.9679 - val\_loss: 0.0593 - val\_accuracy: 0.9810

Epoch 6/15

60000/60000 [=====] - 223s 4ms/step - loss: 0.08

63 - accuracy: 0.9727 - val\_loss: 0.0484 - val\_accuracy: 0.9826

Epoch 7/15

60000/60000 [=====] - 223s 4ms/step - loss: 0.07

43 - accuracy: 0.9766 - val\_loss: 0.0439 - val\_accuracy: 0.9850

Epoch 8/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.06

68 - accuracy: 0.9791 - val\_loss: 0.0404 - val\_accuracy: 0.9865

Epoch 9/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.06

15 - accuracy: 0.9810 - val\_loss: 0.0391 - val\_accuracy: 0.9871

Epoch 10/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.05

83 - accuracy: 0.9818 - val\_loss: 0.0358 - val\_accuracy: 0.9880

Epoch 11/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.05

27 - accuracy: 0.9834 - val\_loss: 0.0355 - val\_accuracy: 0.9880

Epoch 12/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.05

05 - accuracy: 0.9837 - val\_loss: 0.0324 - val\_accuracy: 0.9884

Epoch 13/15

60000/60000 [=====] - 223s 4ms/step - loss: 0.04

79 - accuracy: 0.9847 - val\_loss: 0.0299 - val\_accuracy: 0.9903

Epoch 14/15

60000/60000 [=====] - 223s 4ms/step - loss: 0.04

53 - accuracy: 0.9850 - val\_loss: 0.0287 - val\_accuracy: 0.9901

Epoch 15/15

60000/60000 [=====] - 223s 4ms/step - loss: 0.04

19 - accuracy: 0.9872 - val\_loss: 0.0285 - val\_accuracy: 0.9902

Test loss: 0.028495841736346485

Test accuracy: 0.9901999831199646


CPU times: user 1h 46min 58s, sys: 1min 31s, total: 1h 48min 29s

Wall time: 56min 17s



In [0]:

```
# plotting train and test error  
plot_loss(training, epochs)
```



best epoch is 15. best accuracy is 99.02%

**3 conv layers, kernels = (7x7),(5x5),(3x3), Using tanh**

In [0]:

```
%%time

model = Sequential()
model.add(Conv2D(32, kernel_size=(7, 7),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(128, (3,3), activation='tanh'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.3))

model.add(Flatten())
model.add(Dense(512, activation='tanh'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

training = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_11 (Conv2D)	(None, 18, 18, 64)	51264
max_pooling2d_7 (MaxPooling2D)	(None, 9, 9, 64)	0
dropout_10 (Dropout)	(None, 9, 9, 64)	0
conv2d_12 (Conv2D)	(None, 7, 7, 128)	73856
max_pooling2d_8 (MaxPooling2D)	(None, 3, 3, 128)	0
dropout_11 (Dropout)	(None, 3, 3, 128)	0
flatten_4 (Flatten)	(None, 1152)	0
dense_7 (Dense)	(None, 512)	590336
dropout_12 (Dropout)	(None, 512)	0

dense\_8 (Dense) (None, 10) 5130

=====

Total params: 722,186

Trainable params: 722,186

Non-trainable params: 0

---

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.21

53 - accuracy: 0.9316 - val\_loss: 0.0658 - val\_accuracy: 0.9805

Epoch 2/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.07

55 - accuracy: 0.9766 - val\_loss: 0.0332 - val\_accuracy: 0.9889

Epoch 3/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.06

05 - accuracy: 0.9811 - val\_loss: 0.0320 - val\_accuracy: 0.9889

Epoch 4/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.04

66 - accuracy: 0.9855 - val\_loss: 0.0343 - val\_accuracy: 0.9895

Epoch 5/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.04

48 - accuracy: 0.9862 - val\_loss: 0.0246 - val\_accuracy: 0.9930

Epoch 6/15

60000/60000 [=====] - 225s 4ms/step - loss: 0.03

82 - accuracy: 0.9878 - val\_loss: 0.0295 - val\_accuracy: 0.9909

Epoch 7/15

60000/60000 [=====] - 226s 4ms/step - loss: 0.03

34 - accuracy: 0.9896 - val\_loss: 0.0249 - val\_accuracy: 0.9922

Epoch 8/15

60000/60000 [=====] - 226s 4ms/step - loss: 0.03

02 - accuracy: 0.9901 - val\_loss: 0.0206 - val\_accuracy: 0.9932

Epoch 9/15

60000/60000 [=====] - 227s 4ms/step - loss: 0.02

76 - accuracy: 0.9911 - val\_loss: 0.0213 - val\_accuracy: 0.9937

Epoch 10/15

60000/60000 [=====] - 226s 4ms/step - loss: 0.02

42 - accuracy: 0.9924 - val\_loss: 0.0188 - val\_accuracy: 0.9945

Epoch 11/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.02

41 - accuracy: 0.9921 - val\_loss: 0.0221 - val\_accuracy: 0.9932

Epoch 12/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.02

08 - accuracy: 0.9930 - val\_loss: 0.0211 - val\_accuracy: 0.9935

Epoch 13/15

60000/60000 [=====] - 223s 4ms/step - loss: 0.02

07 - accuracy: 0.9934 - val\_loss: 0.0190 - val\_accuracy: 0.9936

Epoch 14/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.01

72 - accuracy: 0.9941 - val\_loss: 0.0188 - val\_accuracy: 0.9947

Epoch 15/15

60000/60000 [=====] - 224s 4ms/step - loss: 0.01

70 - accuracy: 0.9944 - val\_loss: 0.0228 - val\_accuracy: 0.9925

Test loss: 0.02281436841523082

Test accuracy: 0.9925000071525574


CPU times: user 1h 46min 56s, sys: 1min 34s, total: 1h 48min 30s

Wall time: 56min 25s



In [0]:

```
# plotting train and test error
plot_loss(training, epochs)
```



best epoch is 5. best accuracy is 99.30%

In [4]:

```
pt = PrettyTable()
pt.field_names = ["Layers", "Kernels", "Test Accuracy", "epoch", "activation"]
pt.add_row(["2", "(7x7), (5x5), (3x3)", "99.02", "15", "sigmoid"])
pt.add_row(["3", "(7x7), (5x5), (3x3)", "99.30", "5", "tanh"])
print(pt)
```

Layers	Kernels	Test Accuracy	epoch	activation
2	(7x7), (5x5), (3x3)	99.02	15	sigmoid
3	(7x7), (5x5), (3x3)	99.30	5	tanh

**As number of convolution layers increase test accuracy also increased.**